

Assignment 2 Report

Important Aspects:

Main

The main function first gets the input from the command line. The main function calls `get_cmd_type` function; this function is responsible for parsing the line to their variables, and determining what type of message was given (Type A,B,C, or none of the above). If message is Type B (Create Thread), switch case 1 is responsible for dealing with it. It creates the thread using `pthread_create`, and it inputs the thread into a list of threads. If it is Type C(Terminate), switch case 2 is responsible. It removes the threads with the same Message Type from the thread list, and unassigned alarms for the `alarm_list`, and frees memory. It also terminates the thread using `pthread_cancel`. If the message is Type A(create alarm), Case 3 is responsible. It allocates memory space for an alarm and assigns the struct with the variables from the command. It is placed in the `alarm_list` through function `alarm_insert` by increasing `MessageType`. If it is none of the above, a 'Bad Command' message is shown.

alarm_thread

The alarm thread holds local variables to the thread such as a sublist of messages that have been assigned to it. The alarm thread also pushes and pops a function to stack for when the thread is terminated; `thread_terminate_cleanup` frees memory allocated to the thread's alarms and unlocks mutex for when thread is cancelled. At the start of the while loop, it locks the mutex.

If the alarm list is empty, it goes to `pthread_cond_wait`, which waits for the `pthread_cond_broadcast` in `alarm_insert`; when an alarm is inputted into the list, all threads are made known, so they look in the list. The mutex is unlocked immediately afterwards allowing other threads to be made aware of the alarm.

If the `alarm_list` contains the same message type as the thread, and it has not been assigned, the alarm is assigned to the thread; if not, the thread waits for another alarm in `pthread_cond_wait`. If a suitable alarm is found, it is placed in the threads sublist by increasing time. If the alarm is ready to be outputted, it is outputted. If the alarm is not ready to be outputted, the thread looks into the list again, and assigns them if they are suitable. If a new alarm is assigned, then it is put into the thread sublist. If the new alarm has a shorter time, than it replaces the old alarm, and it is outputted when time is up.

thread_terminate_cleanup

This function is responsible for cleaning up a thread after it has been cancelled by the main thread. The function frees up the memory occupied by its alarms and unlocks the mutex so that other threads have access to the mutex. This function will only be called if the mutex is locked; this function is only called if the main thread calls `pthread_cancel`, and the thread is at `pthread_cond_wait`(mutex is locked after condition arrives), or `sleep(0)`(mutex is locked at this point), since they are the only cancellation points in the threads.

Issues and Solutions:

One problem we faced was the inability to terminate the thread correctly. When 'pthread_cancel(temp_thread->thread_id)' had been called by the main function, the mutex locks were still locked when the thread had been terminated; no other alarm thread was able to be run as a result. To overcome the issue, 'pthread_cleanup_push(thread_terminate_cleanup, (void*)thread_alarm_list)', and 'pthread_cleanup_pop(1)' and been implemented in the alarm thread. Whenever the main thread had called cancel, the thread_terminate_cleanup function had been able to unlock the threads before the thread s terminated.

Another issue we had faced was if a message was entered into the command line before a thread had been created. The time given to the alarm will continue to run even if no thread had been able to service it. To overcome this issue, whenever the alarm does get assigned, the alarm's time is updated, so it is outputted correctly.

An issue we faced was Segmentation Errors that had occurred due to the inability to traverse the alarm_list, and the thread list for when a Type C message had been inputted into command line. This was overcome by debugging the code and printing out messages for where the code was working.

In order for the data to be synchronized, the data has to be locked before accessing it. If the data is not locked, it would lead to inconsistencies if the thread order is not the same as intended. The alarm mutex is locked before the alarm_list is modified and unlocked after modification.

Passing an int to the thread was not straight forward. The value of the message type (was passed in the pthread_create function) changed if multiple threads are created at the same time. To overcome this issue, the message type variable was allocated to memory space, and freed upon the thread receiving it.

If the main thread had multiple input at the same time, it might hog up all the CPU time in-between messages. In order to combat this, the thread was yielded, allowing the other threads to access CPU time.

Quality of Design

The program has modular design, so that functions that are known to work are kept separate from the thread function as well as main. Mutex are used appropriately so that it is not locked if it is already locked, or unlocked if already unlocked, which would have led to undefined behaviour. All memory space that had been allocated, is freed upon termination of the thread.

Testing

All 4 of us had created test cases to see if the output of the program is desired. The test cases vary widely so see if the program hangs, or a segmentation error occurs. Proper design choices were made so that this did not occur.